

# Multiprocessor Systems

- Tightly Coupled vs. Loosely Coupled Systems
  - tightly coupled system generally represent systems which have some degree of sharable memory through which processors can exchange information with normal load / store operations
  - Loosely coupled systems generally represent systems in which each processor has its own private memory and processor to processor information exchange is done via some message passing mechanism like a network interconnect or an external shared channel (FC, IB, SCSI, etc.) bus

# Multiprocessor Systems

- The text presentation in chapters 16 and 17 deals primarily with tightly coupled systems in 2 basic categories:
  - uniform memory access systems (UMA)
  - non-uniform memory access systems (NUMA)
- Distributed systems (discussed beginning in chapter 4) are often referred to as no remote access or NORMA systems

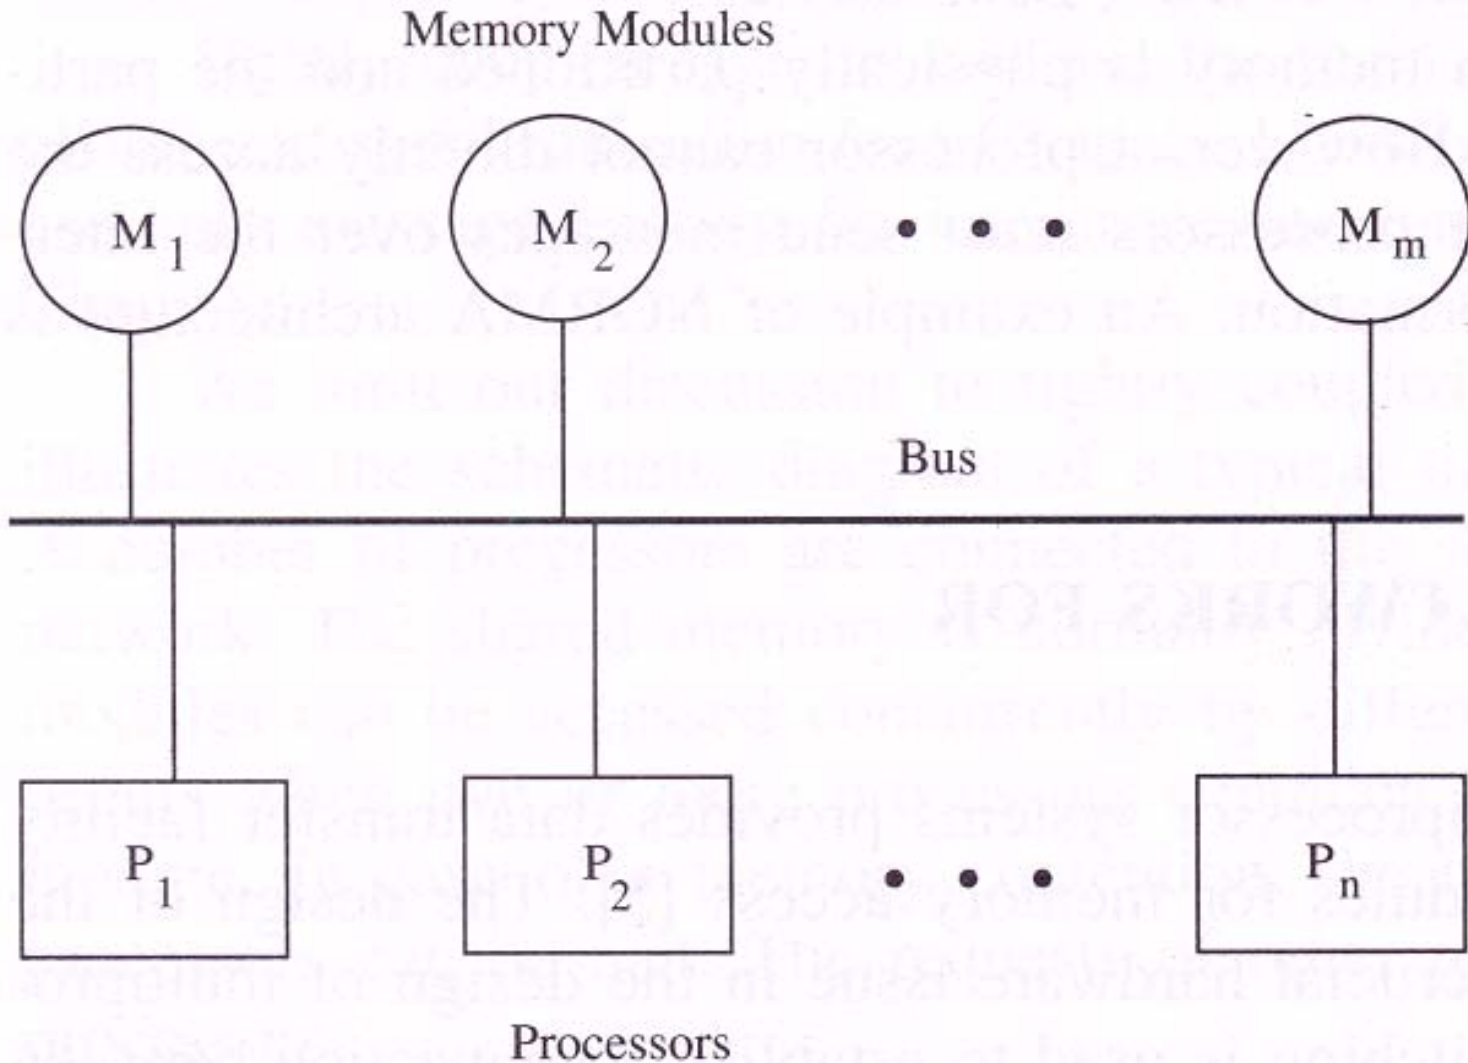
# Multiprocessor Systems

- UMA and NUMA systems provide access to a common set of physical memory addresses using some interconnect strategy
  - a single bus interconnect is often used for UMA systems, but more complex interconnects are needed for scale-up
    - cross-bar switches
    - multistage interconnect networks
  - some form of fabric interconnect is common in NUMA systems
    - far-memory fabric interconnects with various cache coherence attributes

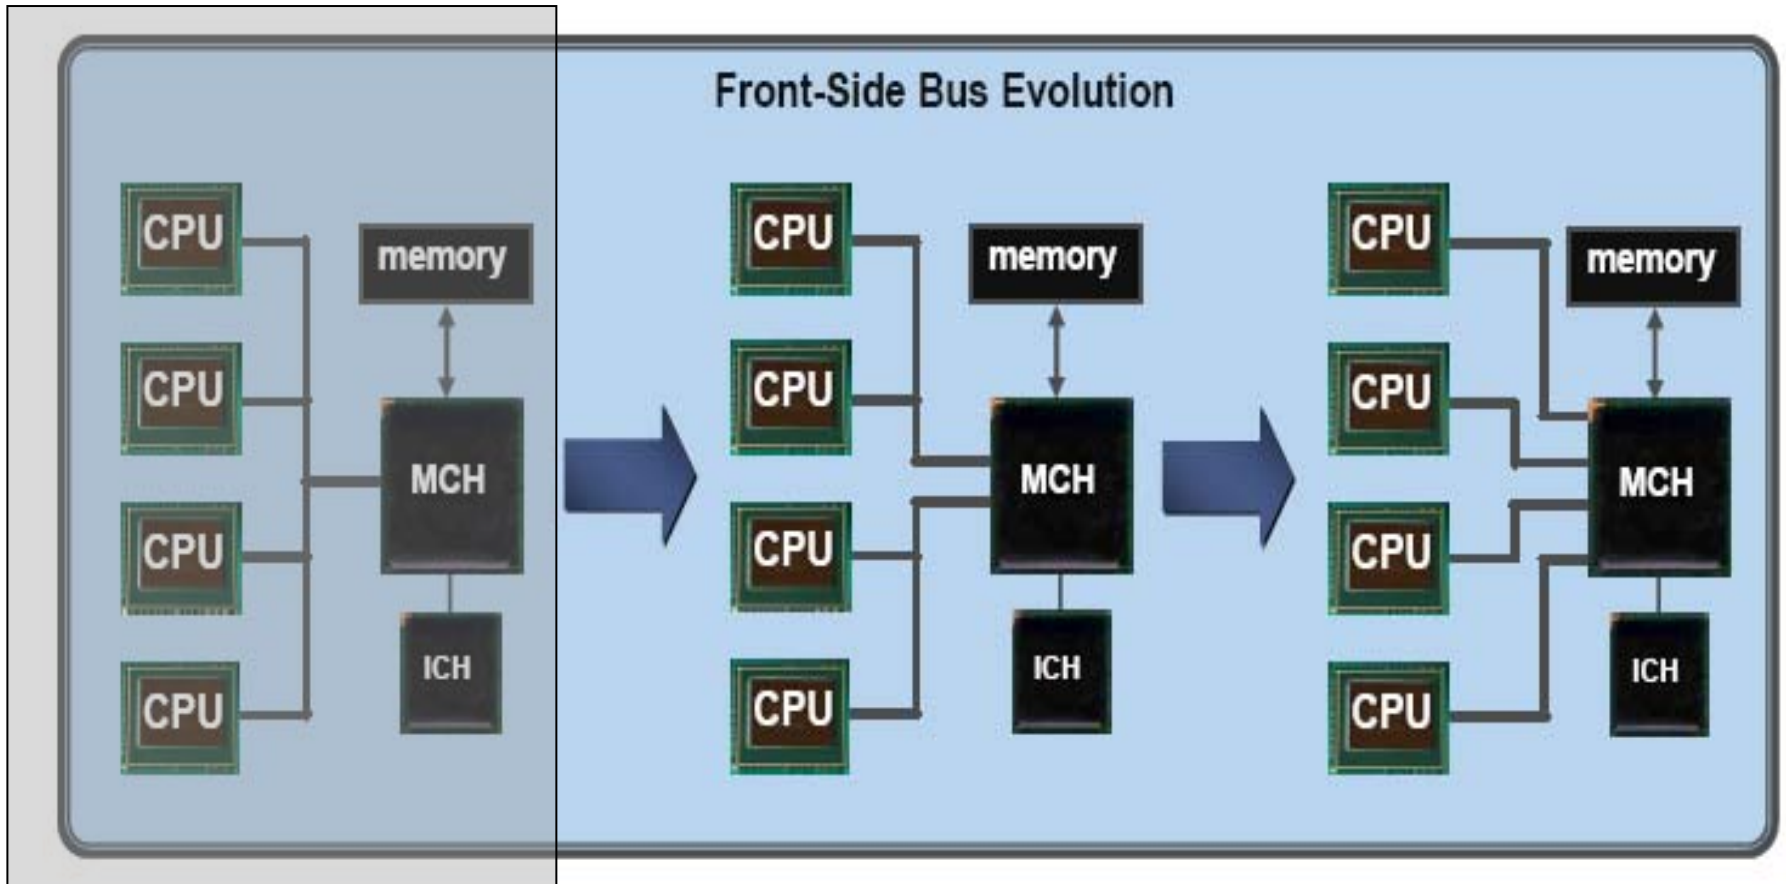
# UMA Bus-Based SMP Architectures

- The simplest multiprocessors are based on a single bus.
  - Two or more CPUs and one or more memory modules all use the same bus for communication.
  - If the bus is busy when a CPU wants to read memory, it must wait.
  - Adding more CPUs results in more waiting.
  - This can be alleviated by having a private cache for each CPU.

# Single Bus Topology



# Single Bus Topology




# UMA Multiprocessors Using Crossbar Switches

- Even with all possible optimizations, the use of a single bus limits the size of a UMA multiprocessor to about 16 CPUs.
  - To go beyond that, a different kind of interconnection network is needed.
  - The simplest circuit for connecting  $n$  CPUs to  $k$  memories is the **crossbar switch**.
    - Crossbar switches have long been used in telephone switches.
    - At each intersection is a **crosspoint** - a switch that can be opened or closed.
    - The crossbar is a **nonblocking** network



# Cross-bar Switch Topology

## Scale-up is $\sim N^2$

Memory Modules

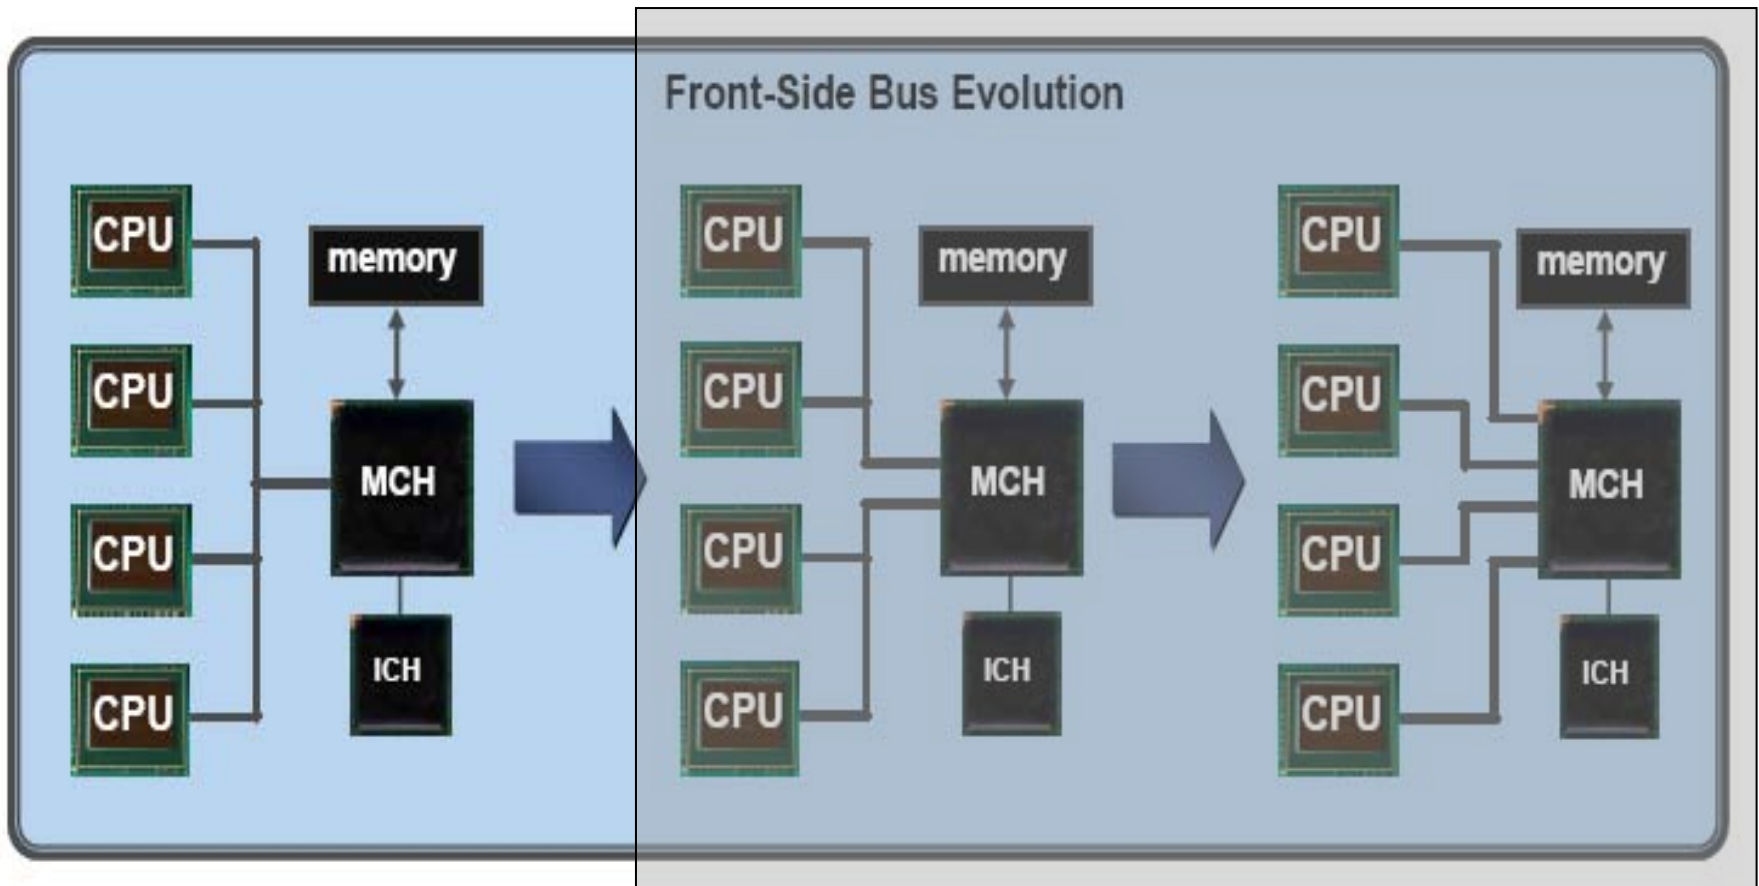


Processors

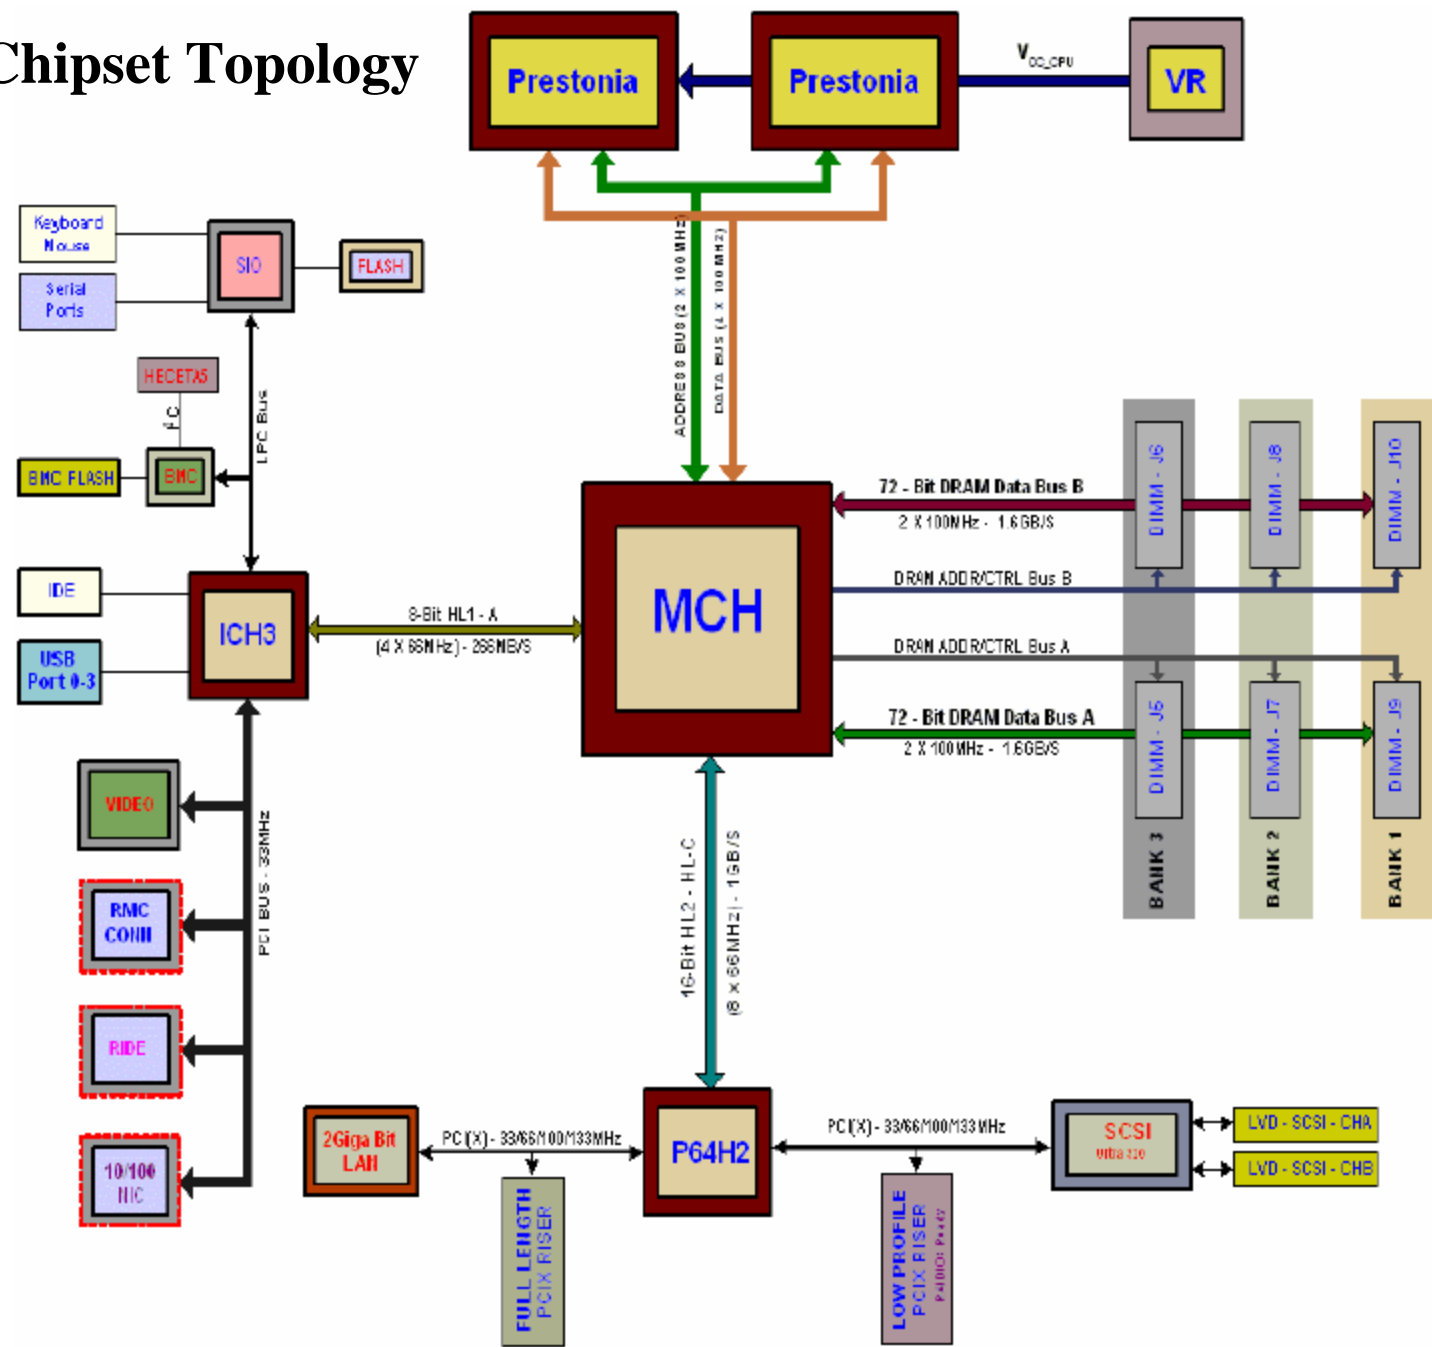




# Crossbar Chipset Topology

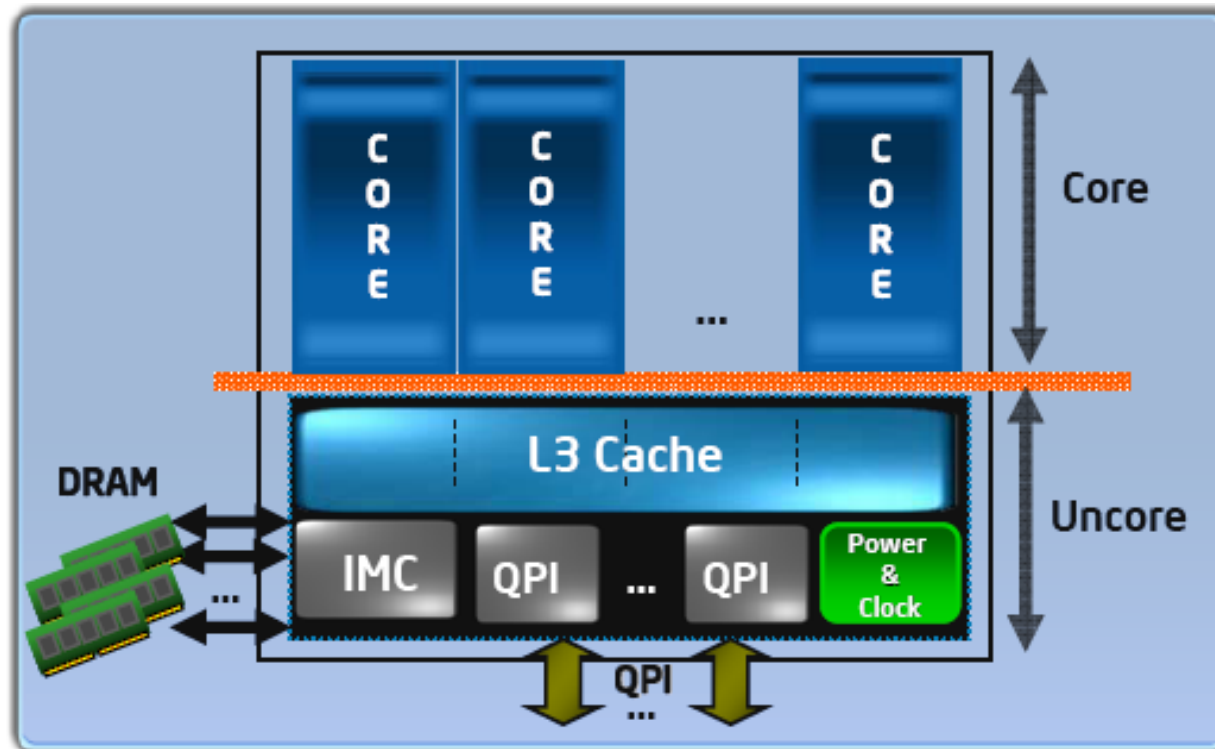


# Crossbar Chipset Topology



# Crossbar On-Die Topology

## Nehalem Core Architecture



QPI: Intel®  
QuickPath  
Interconnect

Differentiation in the "Uncore":

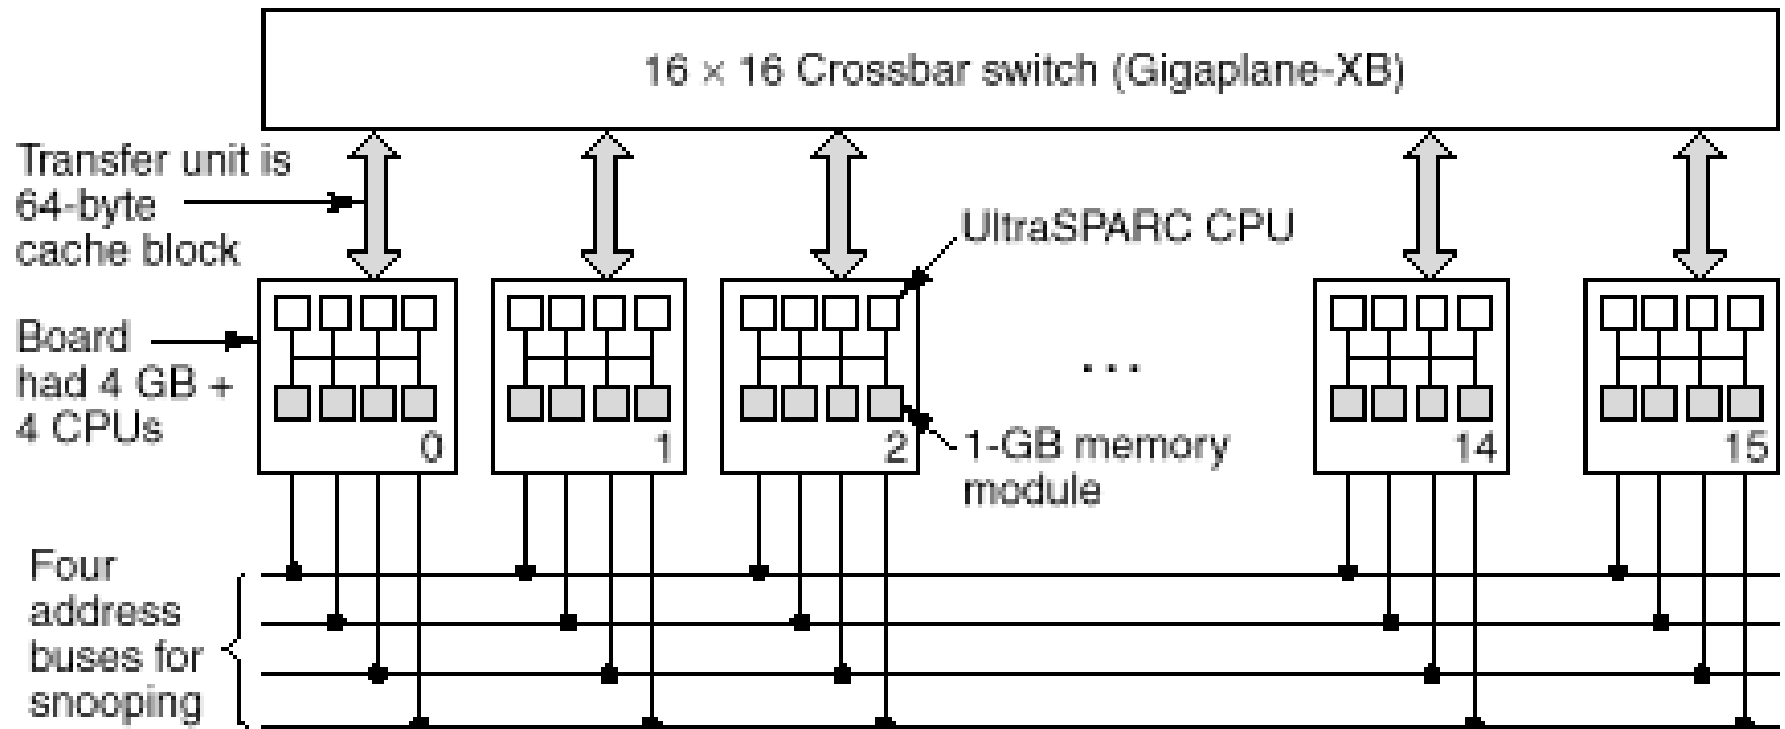


2008 - 2009 Servers & Desktops

# Sun Enterprise 1000

- An example of a UMA multiprocessor based on a crossbar switch is the Sun Enterprise 1000.
  - This system consists of a single cabinet with up to 64 CPUs.
  - The crossbar switch is packaged on a circuit board with eight plug in slots on each side.
  - Each slot can hold up to four UltraSPARC CPUs and 4 GB of RAM.
  - Data is moved between memory and the caches on a 16 X 16 crossbar switch.
  - There are four address buses used for snooping.

# Sun Enterprise 1000 (cont'd)




# UMA Multiprocessors Using Multistage Switching Networks

- In order to go beyond the limits of the Sun Enterprise 1000, we need to have a better interconnection network.
- We can use 2 X 2 switches to build large **multistage switching networks**.
  - One example is the **omega network**.
  - The wiring pattern of the omega network is called the perfect shuffle.
  - The labels of the memory can be used for routing packets in the network.
  - The omega network is a **blocking network**.

# Multistage Interconnect Topology

## Scale-up is $N (\log N)$



# NUMA Multiprocessors


- To scale to more than 100 CPUs, we have to give up uniform memory access time.
- This leads to the idea of **NUMA (NonUniform Memory Access)** multiprocessors.
  - They share a single address space across all the CPUs, but unlike UMA machines local access is faster than remote access.
  - All UMA programs run without change on NUMA machines, but the performance is worse.
    - When the access time to the remote machine is not hidden (by caching) the system is called **NC-  
NUMA**.



# NUMA Multiprocessors (cont'd)

- When coherent caches are present, the system is called **CC-  
NUMA**.
- It is also sometimes known as **hardware DSM** since it is basically the same as software distributed shared memory but implemented by the hardware using a small page size.
- One of the first NC-NUMA machines was the Carnegie Mellon Cm\*.
  - This system was implemented with LSI-11 CPUs (the LSI-11 was a single-chip version of the DEC PDP-11).
  - A program running out of remote memory took ten times as long as one using local memory.
  - Note that there is no caching in this type of system so there is no need for cache coherence protocols

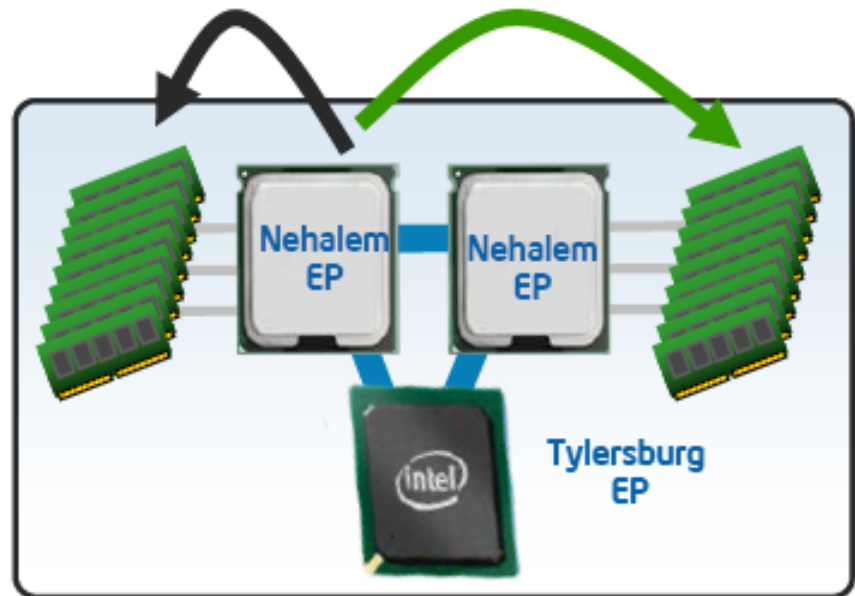
In a full NUMA system memory and peripheral space is visible to any processor on any node



# NUMA On-Die Topology

## Nehalem Core Architecture

- FSB architecture
  - All memory in one location
- Starting with Nehalem
  - Memory located in multiple places
- Latency to memory dependent on location
- Local memory
  - Highest BW
  - Lowest latency
- Remote Memory
  - Higher latency





***Ensure software is NUMA-optimized for best performance***

# NUMA QPI Support


## Nehalem Core Architecture

- Nehalem introduces new Intel® QuickPath Interconnect (QPI)
- **High bandwidth, low latency** point to point interconnect
- Up to 6.4 GT/sec initially
  - 6.4 GT/sec -> 12.8 GB/sec
  - Bi-directional link -> 25.6 GB/sec per link
  - Future implementations at even higher speeds
- Highly **scalable** for systems with varying # of sockets




# Local Memory Access

- CPU0 requests cache line X, not present in any CPU0 cache
  - CPU0 requests data from its DRAM
  - CPU0 snoops CPU1 to check if data is present
- Step 2:
  - DRAM returns data
  - CPU1 returns snoop response
- Local memory latency is the maximum latency of the two responses
- Nehalem optimized to keep key latencies close to each other



# Remote Memory Access

- CPU0 requests cache line X, not present in any CPU0 cache
  - CPU0 requests data from CPU1
  - Request sent over QPI to CPU1
  - CPU1's IMC makes request to its DRAM
  - CPU1 snoops internal caches
  - Data returned to CPU0 over QPI
- Remote memory latency a function of having a low latency interconnect



# Multiprocessor Operating Systems

- Common software architectures of multiprocessor systems:
  - Separate supervisor configuration
    - Common in clustered systems
    - May only share limited resources
  - Master-Slave configuration
    - One CPU runs the OS, others run only applications
    - OS CPU may be a bottleneck, may fail
  - Symmetric configuration (SMP)
    - One OS runs everywhere
    - Each processor can do all (most) operations


# Multiprocessor Operating Systems

- SMP systems are most popular today
  - Clustered systems are generally a collection of SMP systems that share a set of distributed services
- OS issues to consider
  - Execution units (threads)
  - Synchronization
  - CPU scheduling
  - Memory management
  - Reliability and fault tolerance



# Multiprocessor Operating Systems

- Threads (execution units)
  - Address space utilization
  - Platform implementation
    - User level threads ( $M \times 1$  or  $M \times N$ , Tru64Unix, HP-UX)
      - Efficient
      - Complex
      - Course grain control
    - Kernel level threads ( $1 \times 1$ , Linux, Windows)
      - Expensive
      - Less complex
      - Fine grain control



**Process 2 is equivalent to a pure ULT approach**

**Process 4 is equivalent to a pure KLT approach**

**We can specify a different degree of parallelism (process 3 and 5)**

# Multiprocessor Operating Systems

- Synchronization issues
  - Interrupt disable no longer sufficient
  - Spin locks required
    - Software solutions like Peterson's algorithm are required when hardware platform only offers simple memory interlock
    - Hardware assist needed for efficient synchronization solutions
      - Test-and-set type instructions
        - » Intel XCHG instruction
        - » Motorola 88110 XMEM instruction
      - Bus lockable instructions
        - » Intel CMPXCHG8B instruction

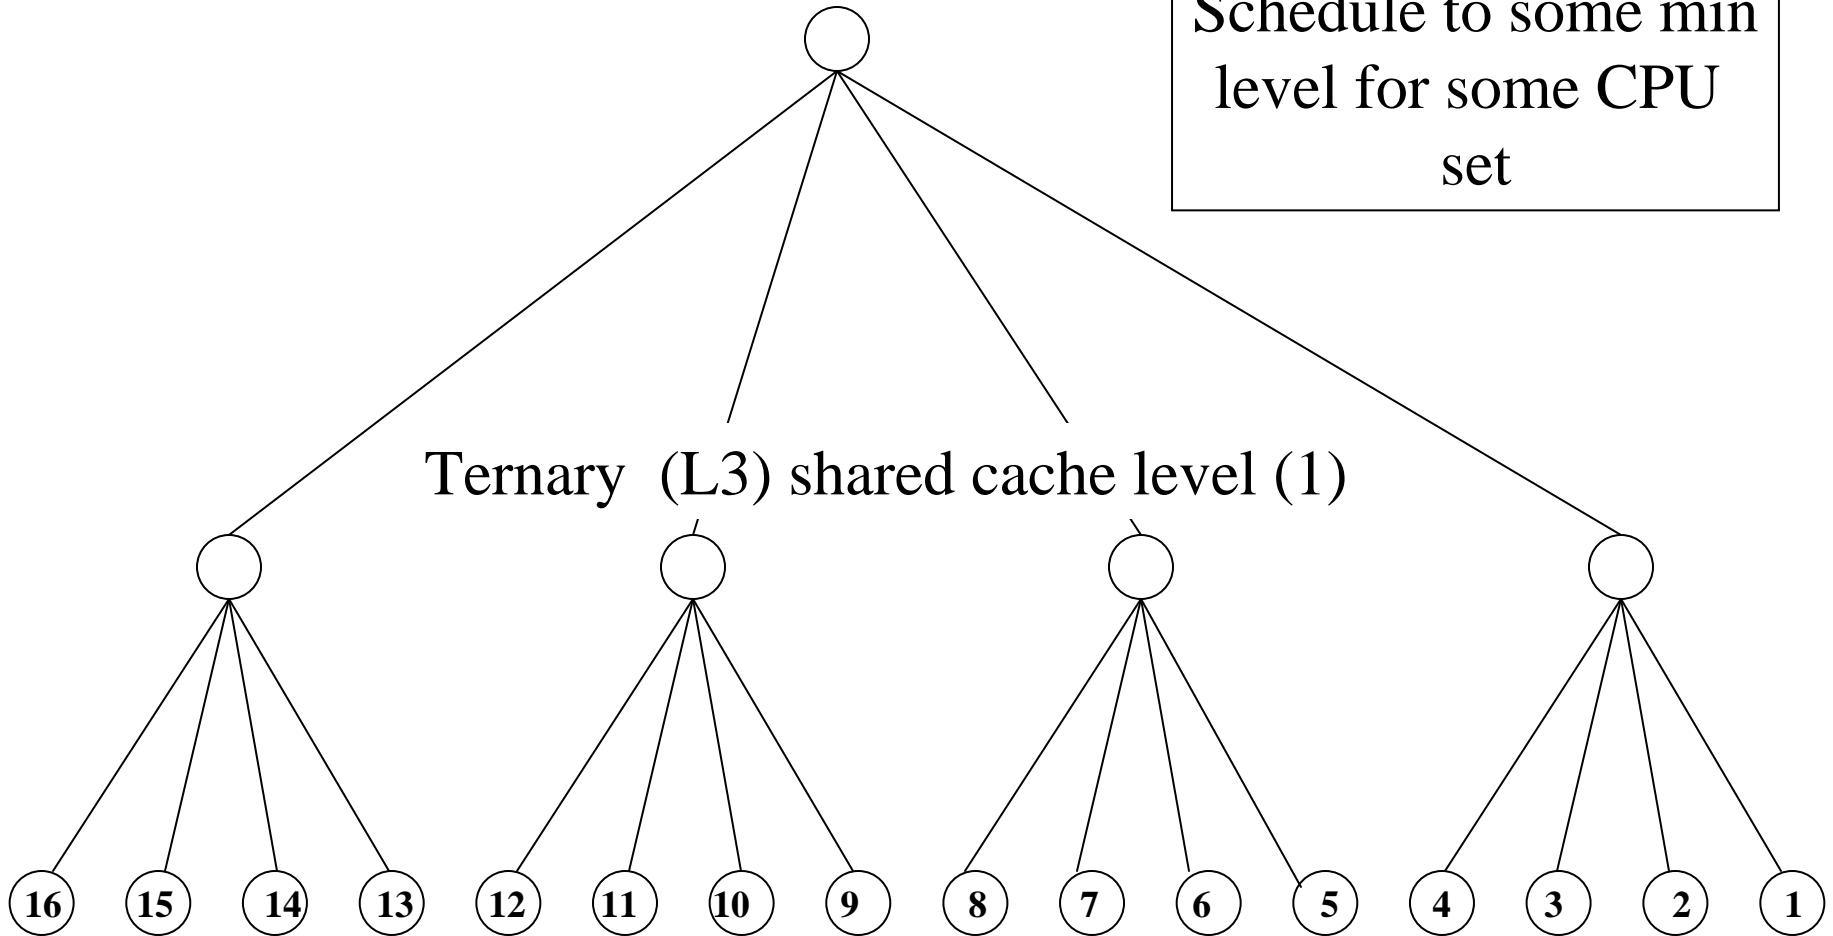
# Multiprocessor Operating Systems

- Processor scheduling issues
  - Schedule at the process or thread level ?
  - Which processors can an executable entity be scheduled to ?
    - Cache memory hierarchies play a major role
      - Affinity scheduling and cache footprint
    - Can the OS make good decisions without application hints ?
      - Applications understand how threads use memory, the OS does not
    - What type of scheduling policies will the OS provide ?
      - Time sharing, soft/hard real time, etc.

Main memory level (2)

Schedule to some min  
level for some CPU  
set

Ternary (L3) shared cache level (1)



CPU and dedicated cache (L1,L2) level (0)

# Multiprocessor Operating Systems

- Memory management issues
  - Physical space deployment (UMA, NUMA)
  - Address space organization
  - Types of memory objects
    - Text, data, stack, heap, memory mapped files, shared memory segments, etc.
    - Shared vs private objects
    - Anonymous vs file objects
      - Swap space allocation strategies
  - Paging strategies and free memory list(s) configurations
  - Kernel data structure locations

# Reliability and Fault Tolerance Issues

- Operating systems must keep complex systems alive
  - Simple panics no longer make sense
- OS must be proactive in keeping the system up
  - Faulty components must be detected and isolated to the degree enabled by HW
  - The system must map its way around failed HW and continue to run
  - To the extent that the HW supports hot repair, the OS must provide recovery mechanisms