

Process Management (Cont'd)

Get your own PID

Get your parent's PID

Get your own process group ID

SYNOPSIS

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t getpid(void);
```

```
pid_t getppid(void);
```

```
pid_t getpgrp(void);
```

returns: PID of interest, cannot fail

EXAMPLE:

```
printf("My PID is %d\n", getpid());  
printf("My parent's PID is %d\n", getppid());  
printf("My process group ID is %d\n",  
       getpgrp());
```

Process Management (Cont'd)

Get real and effective UIDs and GIDs

SYNOPSIS

```
#include <unistd.h>
#include <sys/types.h>
```

```
uid_t    getuid ()
```

```
gid_t    getgid ()
```

```
uid_t    geteuid ()
```

```
gid_t    getegid ()
```

returns: ID of interest, cannot fail

EXAMPLE:

```
printf("My RUID is %d\n", getuid());
printf("My RGID is %d\n", getgid());
printf("My EUID is %d\n", geteuid());
printf("My EGID is %d\n", getegid());
```

Process Management (Cont'd)

Set real and effective UIDs and GIDs

SYNOPSIS

```
#include <unistd.h>
#include <sys/types.h>
```

```
uid_t    setuid (uid)
uid_t    uid;
```

```
uid_t    setgid (gid)
gid_t    gid;
```

returns: 0 on success or -1

EXAMPLE:

```
int rgid, ruid;

printf("My RUID is %d\n", getuid());
printf("My RGID is %d\n", getgid());
printf("My EUID is %d\n", geteuid());
printf("My EGID is %d\n", getegid());

if(setuid(ruid) == -1 | setgid(rgid) == -1){
    perror("setting IDs back failed: ");
    exit(1)
}

printf("My EUID is now %d\n", geteuid());
printf("My EGID is now %d\n", getegid());
```

Process Management (Cont'd)

Set session and group IDs for job control

SYNOPSIS

```
#include <sys/types.h>
#include <unistd.h>
```

```
pid_t  setsid ()
```

```
int setpgid (pid, pgid)
pid_t pid, pgid;
```

where:

pid The process id of the process whose
 process group id is to be changed.

pgid The new process group id.

returns: returns: 0 on success or -1

EXAMPLE:

```
int my_pid;

if((my_pid = getpid()) != getpgrp()){
    if(setpgid((pid_t)0, my_pid) == -1){
        perror("group leader attempt failed: ");
        exit(1);
    }
}
```

EXAMPLE:

```
if(setsid() == -1){
    perror("cannot become session leader: ");
    exit(1);
}
```

Process Management (Cont'd)

Process priority adjustment

SYNOPSIS

```
#include <unistd.h>
```

```
int  nice (incr)
```

```
int  incr;
```

where:

incr A positive or negative value that
 is to be added to the calling process's
 priority

returns: New nice value on success or -1 on
 error. Always clear errno before call

DESCRIPTION

The value of incr is added to the priority of the calling process. A more positive priority value results in a lower level of service from the CPU.

If the new priority would be greater than 19, the process's priority is set to 19. If the new priority would be less than -20, the process's priority is set to -20.

Process Management (Cont'd)

`nice()` cont'd

EXAMPLE:

```
int new_pri;

printf("current nice value is %d\n", nice(0));
if ((new_pri = nice(+8)) == -1){
    perror("nice change failed: ");
    exit(1);
}
printf("new nice value is %d\n", new_pri);

printf("current nice value is %d\n", nice(0));
if ((new_pri = nice(-8)) == -1){
    perror("nice privileged change failed: ");
    exit(1);
}
printf("new privileged nice value is %d\n", new_pri);
```

Process Management (Cont'd)

Process priority adjustment BSD style

SYNOPSIS

```
#include <sys/resource.h>
```

```
int  getpriority (which, who)
```

```
int  which;
```

```
int  who;
```

```
int  setpriority (which, who, prio)
```

```
int  which;
```

```
int  who;
```

```
int  prio;
```

where:

which How the argument who is to be interpreted
 in identifying one or more processes
 whose priorities will be set:
 PRIO_PROCESS, PRIO_PGRP, or PRIO_USER

who Identifier of one or more processes
 whose priorities will be set: a process
 ID, a process group ID, or user ID,
 depending on the value of which

prio The new priority value (range -20 to 20)

returns: getpriority -- new nice value on success
 setpriority -- 0 on success
 both -- -1 on error

Process Management (Cont'd)

getpriority(), setpriority() cont'd

EXAMPLE:

```
int max_uid_old_pri, new_pri;

if((max_uid_old_pri =
    getpriority(PRIO_USER, getuid())) == -1){
    perror("getpriority failed: ");
    exit(1);
}

new_pri = max_uid_old_pri + 10;
if(setpriority(PRIO_USER, getuid(), new_pri) == -1){
    perror("set un-privileged priority failed: ");
    exit(1);
}

new_pri = max_uid_old_pri - 10;
if(setpriority(PRIO_USER, getuid(), new_pri) == -1){
    perror("set privileged priority failed: ");
    exit(1);
}
```


Signal Management

UNIX systems employ an asynchronous process notification mechanism known as the **signal** facility. Signals often cause portability problems since there are several similar but different signal implementations. The common implementations include:

- USL System VR3 (22 defined signals)
- USL System VR4 (32 defined signals)
- BSD 4.3 (32 defined signals)
- OSF/1 (32 defined signals)
- DGUX (64 possible, 38 defined)

Even where the number of signals is the same, the actual signals in the set are often slightly different.

Signal Management (Cont'd)

Signals can be used as a means of notifying a process of some event in an asynchronous way. In essence, a signal is sent from a process to a process, and the sender and receiver may be the same process.

Signals are organized into two main groups, those called *synchronous* are always delivered as a result of a process run-time exception, and the offending process delivers the signal to itself during exception processing. Synchronous signals include **SIGSEGV**, **SIGPIPE** and **SIGFPE** , among others.

The so-called *asynchronous* signals are delivered to a process by another process, and are in no direct way tied to any specific behavior of the target process. Asynchronous signals include **SIGINT**, **SIGTERM** and **SIGKILL** , among others.

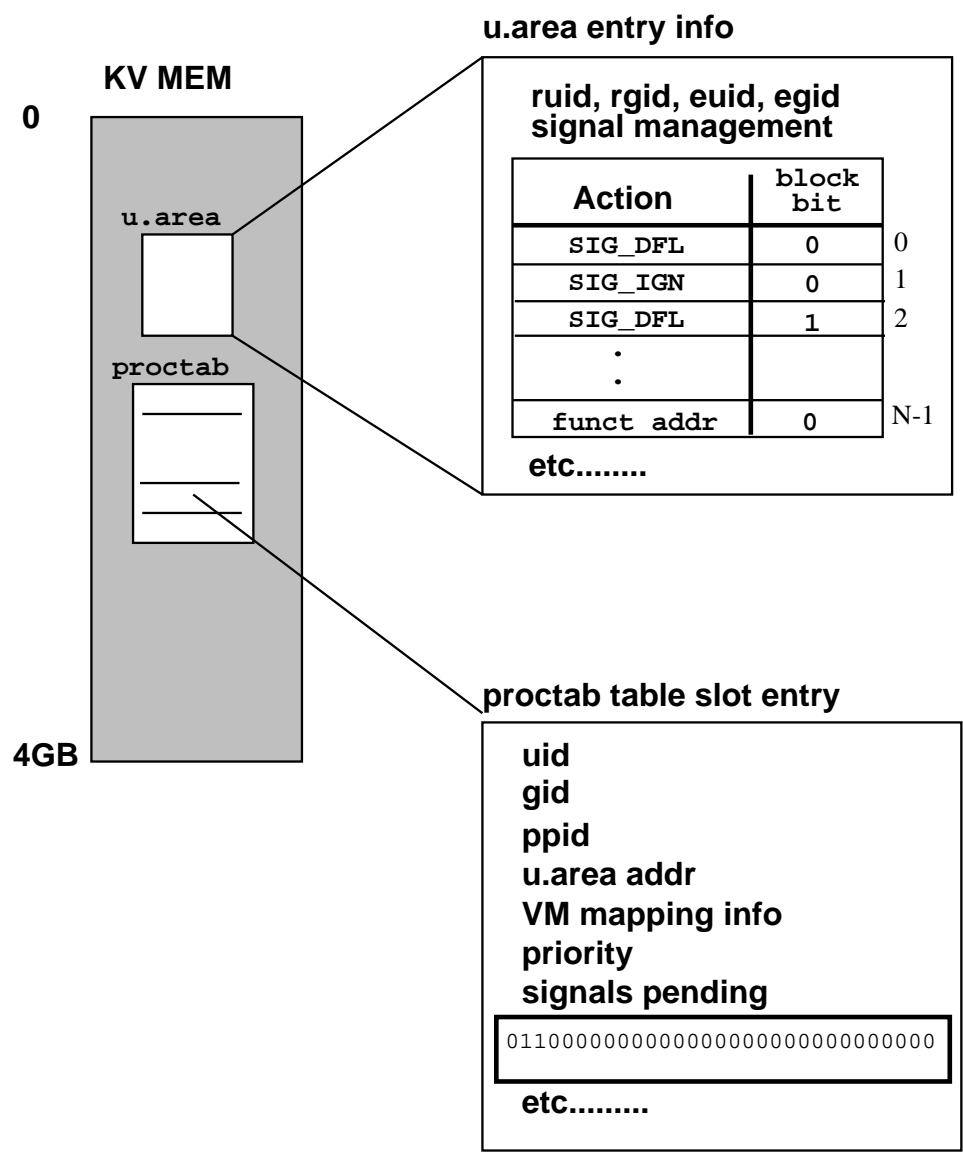
Signal Management (Cont'd)

Each process keeps a *signal action table* in the *u.area* and a *pending signal vector* in its **proctab** entry. When a signal is launched from a sending process to a target process(es), a bit is turned on in the pending signal vector of the proctab entry of the target process.

Before a process is allowed to transition from *kernel running* to *user running* the process must check its pending signal vector for any outstanding signals, and must manage such signals before it can return to its place in user space.

Just exactly how the detecting process will handle a specific signal depends on what information the process has in its *u.area* signal action table in the table slot which corresponds to the outstanding signal.

Signal Management (Cont'd)



Signal Management (Cont'd)

A process can arrange to handle arriving signals in one of three basic ways:

- The process can set the signal for a signal defined default action with **SIG_DFL**
- The process can set the signal to be discarded upon arrival with **SIG_IGN**
- The process can install a user supplied function address pointing to a function to be invoked should the signal arrive

Most signal defined default actions force the process to terminate by calling `exit` upon signal detection.

User defined signal handling functions must be of type **void** and will be passed a single integer parameter when invoked, which is the signal number which caused the invocation.

If the *block bit* is set to **1** for a given signal and the signal should arrive to a process, the process will not attempt to handle the signal until the corresponding block bit is set to **0**. If additional instances of the signal should arrive while the block bit is set, such instances will be queued for future disposition.

Signal Management (Cont'd)

The following signals are defined for USL SVR4

SIGHUP	1	hangup
SIGINT	2	interrupt
SIGQUIT	3*	quit
SIGILL	4*	illegal instruction
SIGTRAP	5*	trace trap
SIGABRT	6*	abort
SIGEMT	7*	EMT instruction
SIGFPE	8*	floating point exception
SIGKILL	9	kill (cannot be caught, blocked, or ignored)
SIGBUS	10*	bus error
SIGSEGV	11*	segmentation violation
SIGSYS	12*	bad argument to system call
SIGPIPE	13	write on a pipe, no open read
SIGALRM	14	alarm clock
SIGTERM	15	software termination signal
SIGUSR1	16	user signal 1
SIGUSR2	17	user signal 2
SIGCHLD	18@	child status changed
SIGPWR	19@	power fail/restart
SIGWINCH	20@	signal window size change
SIGURG	21@	urgent socket condition
SIGPOLL	22	pollable event
SIGSTOP	23 +	stop (cannot be caught, blocked, or ignored)

Signal Management (Cont'd)

USL SVR4 Signals cont'd

SIGTSTP	24 +	stop signal from keyboard
SIGCONT	25@	continue after stop
SIGTTIN	26 +	background read attempted
SIGTTOU	27 +	background write attempted
SIGVTALRM	28	virtual time alarm
SIGPROF	29	profiling timer alarm
SIGXCPU	30*	CPU time limit exceeded
SIGXFSZ	31*	file size limit exceeded

where: no indication means default termination
* means termination with a core dump
|+ means specific job control action
@ means default ignore action

Signal Management (Cont'd)

Examine or change signal action

SYNOPSIS

```
#include <signal.h>

int      sigaction (sig, act, oact)
int      sig;
const struct sigaction *act;
struct sigaction *oact;
```

where:

sig	A signal number.
act	NULL, or a new action to be installed for sig.
oact	NULL, or the current action associated with sig. If act is not NULL and the call is successful, oact will be replaced by act.

```
struct sigaction{
    void (*)(()) sa_handler; /* SIG_DFL, SIG_IGN, or
                                pointer to a function. */
    sigset_t      sa_mask;    /* Additional set of signals
                                to be blocked during execution
                                of signal-catching function. */
    int           sa_flags ;  /* Special flags to affect
                                behavior of signal. */
};
```


Signal Management (Cont'd)

Use the `sa_flags` field to modify the delivery of the specified signal. The values you can specify, defined in the header `<sys/signal.h>`, are listed below:

<code>SA_ONSTACK</code>	If set and the signal is caught, and an alternate signal stack has been declared, the signal is delivered to the calling process using the alternate stack. Otherwise, the signal is delivered on the same stack as the main program.
<code>SA_RESETHAND</code>	If set and the signal is caught, the action of the signal is reset to <code>SIG_DFL</code> . (Note: <code>SIGKILL</code> , <code>SIGTRAP</code> , and <code>SIGPWR</code> cannot be automatically reset when delivered. With these signals, setting <code>SA_RESETHAND</code> has no effect.)
<code>SA_NODEFER</code>	If set and the signal is caught, <code>sig</code> will not be automatically blocked while the handler is active.
<code>SA_RESTART</code>	If set and the signal is caught, and if the system call is restartable, the kernel will restart the system call on behalf of the caller after a signal handler completes processing some signal that has interrupted the call. If this flag is not set, a system call that is interrupted will return <code>EINTR</code> . A non-restartable system call that is interrupted will return <code>EINTR</code> regardless of this flag.

Signal Management (Cont'd)

SA_SIGINFO	If this flag is not set and the signal is caught, sig is passed as the only argument to the signal handling function. If this flag is set and the signal is caught, two additional arguments will be passed to the signal handling function. If the second argument is not equal to NULL, it will point to an object of type <code>siginfo_t</code> , which will explain the reason the signal was generated (see <code>siginfo.h</code>). The third argument will point to an object of type <code>ucontext_t</code> , which will describe the receiving process' context at the time it received the signal (see <code>ucontext.h</code>).
SA_NOCLDWAIT	If set and sig equals SIGCHLD, the system will clean up after the calling process's dead children. If the calling process subsequently calls <code>wait()</code> , it will block until all of its processes terminate and then return a value of -1 with <code>errno</code> set to ECHILD.
SA_NOCLDSTOP	If set and sig equals SIGCHLD, sig will not be sent to the calling process when its child processes stop.

Signal Management (Cont'd)

sigaction() cont'd

EXAMPLE:

```
    struct sigaction  new, old;
    sigset_t  mask_sigs;
    int  i, nsigs;
    int sigs[] = {SIGHUP,SIGINT,SIGQUIT,
                  SIGTERM, SIGXFSZ};

    void  handler(int signum);

        .
        .

    nsigs = sizeof(sigs)/sizeof(int)
    sigemptyset(&mask_sigs);
    for(i=0; i< nsigs; i++)
        sigaddset(&mask_sigs, sigs[i]);
    for(i=0; i< nsigs; i++){
        new.sa_handler = handler;
        new.sa_mask = mask_sigs;
        new.sa_flags = SA_RESTART;
        if(sigaction(sigs[i], &new, &old) == -1){
            perror("can't set signals: ");
            exit(1);
        }
        printf("signal # %d previous action was %x\n",
               sigs[i],  old.sa_handler);
    }
}  /* close main  */
```

Signal Management (Cont'd)

sigaction() cont'd

```

        .
        .
        .
        .
void handler(int signum)
{
    switch(signum){
        case SIGHUP: printf("SIGHUP caught\n")
                      /* clean up environment */
                      printf("going down on SIGHUP\n");
                      exit(2);

        case SIGINT: /* interrupt code */
                      .
                      .
                      .
        default:      /* unexpected sig ?? */
    }
}
```

Signal Management (Cont'd)

Manipulate sets of signals (from libc)

SYNOPSIS

```
#include <signal.h>
```

```
int sigemptyset (sigset_t *set);
int sigfillset (sigset_t *set);
int sigaddset (sigset_t *set, int signo);
int sigdelset (sigset_t *set, int signo);
int sigismember (sigset_t *set, int signo);
```

```
returns:      sigismember -- 1 if true, 0 if false
              all others  -- 0 on success or -1
```

DESCRIPTION

`sigemptyset` initializes the set pointed to by `set` to exclude all signals defined by the system.

`sigfillset` initializes the set pointed to by `set` to include all signals defined by the system.

`sigaddset` adds the individual signal specified by the value of `signo` to the set pointed to by `set`.

`sigdelset` deletes the individual signal specified by the value of `signo` from the set pointed to by `set`.

`sigismember` checks whether the signal specified by the value of `signo` is a member of the set pointed to by `set`.

Any object of type `sigset_t` must be initialized by applying either `sigemptyset` or `sigfillset` before applying any other operation.

Signal Management (Cont'd)

Examine and change the block bits

SYNOPSIS

```
#include <signal.h>

int sigprocmask (how, set, oset)
int how;
const sigset_t *set;
sigset_t      *oset;
```

where:

how The manner in which the current set
 of blocked signals is changed.

set NULL, or the signal set used to change
 the current set of blocked signals.

oset NULL, or the current set of blocked
 signals.

POSSIBLE "how" VALUES

SIG_BLOCK	The resulting set shall be the union of the current set and the signal set pointed to by the argument set.
SIG_UNBLOCK	The resulting set shall be the intersection of the current set and the complement of the signal set pointed to by the argument set.
SIG_SETMASK	The resulting set shall be the signal set pointed to by the argument set.

Signal Management (Cont'd)

Send a signal to a process(es)

SYNOPSIS

```
#include <sys/types.h>
#include <signal.h>
```

```
int kill (pid, sig)
pid_t pid;
int sig;
```

where:

pid	An integer (positive, negative, or zero) indicating a process or a group of processes to be sent the signal
sig	A signal number that is either one from the list given in <signal.h> or zero

If pid is greater than zero, sig shall be sent to the process whose process ID is equal to pid.

If pid is zero, sig shall be sent to all processes (excluding an implementation-defined set of system processes) whose process group ID is equal to the process group ID of the sender, and for which the process has permission to send a signal.

If pid is negative, but not -1, sig shall be sent to all processes whose process group ID is equal to the absolute value of pid, and for which the process has permission to send a signal.

If the value of pid causes sig to be generated for the sending process, and if sig is not blocked, either sig or at least one pending unblocked signal shall be delivered to the sending process before the kill() function returns.