

Signal Management (Cont'd)

Check for pending, blocked signals

SYNOPSIS

```
#include <signal.h>

int sigpending (set)
sigset_t *set;
```

where:

set A structure to which the list of
 signals are to be written

returns: 0 on success or -1 (generally no
 error is possible, but
 implementation defined)

DESCRIPTION

The sigpending() function shall store the set of signals that are blocked from delivery and pending for the calling process, in the space pointed to by the argument set.

Signal Management (Cont'd)

EXAMPLE:

```
sigset_t    pending_sigs, take_sig_mask;
struct sigaction new, old;

if(sigpending (&pending_sigs) == -1){
    perror("sigpending failed: ");
    exit(1);
}

if(sigismember (&pending_sigs, SIGINT)){
    /** SIGINT has been delivered but the  **/
    /** block bit for SIGINT is set so no  **/
    /** action has been taken yet...may    **/
    /** want to install a specific handler **/
    /** and turn off the block bit to      **/
    /** handle the signal now.....       **/

    sigemptyset (&take_sig_mask);
    sigaddset (&take_sig_mask, SIGINT);
    sigaction(SIGINT, &new, &old);
    sigprocmask(SIG_UNBLOCK, &take_sig_mask, NULL);
    pause();
    sigprocmask(SIG_BLOCK, &take_sig_mask, NULL);
    .
    .
    .
}
```

Signal Management (Cont'd)

Wait for a signal to arrive

SYNOPSIS

```
int pause(void)
```

SYNOPSIS

```
#include <signal.h>
```

```
int sigsuspend (sigmask)
sigset_t *sigmask;
```

where:

sigmask A structure containing a set of signals, which temporarily replaces the current block mask and allows the process to sleep until some signal, other than any in the new mask, arrives

returns: -1 with errno EINTR (cannot succeed)
 BUT the original signal mask IS
 restored upon return with sigsuspend()

Signal Management (Cont'd)

`sigsuspend()`, `pause()` cont'd

EXAMPLE:

```
void handler(int signum)
{ return; }

sigset_t    suspended_blocked_sigs, take_sig_mask;
struct sigaction new, old;

sigfillset (suspended_blocked_sigs);
sigfillset (take_sig_mask);
sigdelset(suspended_blocked_sigs, SIGALRM);
new.sa_handler = handler;
new.sa_mask = take_sig_mask;
new.sa_flags = 0;
if(sigaction(SIGALRM, &new, &old) == -1){
    perror("can't set signals: ");
    exit(1);
}
alarm(5);
sigsuspend(suspended_blocked_sigs);
/**** when alarm clock rings, reset signals ****/
/**** how would use of pause() be different ****/

if(sigaction(SIGALRM, &old, NULL) == -1){
    perror("can't set signals: ");
    exit(2);
}
```

Signal Management (Cont'd)

Set the real-time alarm clock

SYNOPSIS

```
#include <unistd.h>
```

```
unsigned int alarm (sec)
```

```
unsigned int sec;
```

where:

sec The number of wall-clock seconds to wait before sending SIGALRM to the caller. If 0, cancel the alarm

returns: time remaining on the clock or 0

EXAMPLE:

```
void handler(int signum)
{ /** handle alarm timeout **/ }
int time_in_code;

if(sigaction(SIGALRM, &new, &old) == -1){
    perror("can't set signals: ");
    exit(1);
}
alarm(5);
/** do code which may encounter black hole **/
time_in_code = 5 - alarm(0);
```

Signal Management (Cont'd)

Set the real-time, virtual or profile alarm clock

SYNOPSIS

```
#include <sys/time.h>
```

```
int setitimer(int which,  
              struct itimerval *value,  
              struct itimerval *ovalue);
```

where:

which	ITIMER_REAL, ITIMER_VIRTUAL, or ITIMER_PROF
value	Name of pointer to structure for storing timer value
ovalue	Name of pointer to structure for storing old timer value

DESCRIPTION

The system provides each process with three interval timers, defined in `sys/time.h`. The `getitimer` call stores the current value of the timer specified by `which` into the structure pointed to by `value`. The `setitimer` call sets the value of the timer specified by `which` to the value specified in the structure pointed to by `value`, and if `ovalue` is not `NULL`, stores the previous value of the timer in the structure pointed to by `ovalue`.

returns: 0 on success, or -1

Signal Management (Cont'd)

setitimer() cont'd

The three timers are:

ITIMER_REAL	Decrements in real time. Delivers SIGALRM signal.
ITIMER_VIRTUAL	Decrements in process virtual time. Delivers SIGVTALRM signal.
ITIMER_PROF	Decrements both in process virtual time and during system call execution. Delivers SIGPROF signal.

```
struct itimerval{
    struct timeval it_interval; /* timer interval */
    struct timeval it_value;    /* current value */
};
```

```
struct timeval{
    long    tv_sec;              /* seconds          */
    long    tv_usec;            /* micro seconds   */
};
```

Signal Management (Cont'd)

setitimer() cont'd

EXAMPLE:

```
void handler(int signum)
{ /** handle virtual alarm timeout */ }

struct sigaction new, old;
struct itimerval *value;
struct itimerval *ovalue;

new.sa_handler = handler;
new.sa_mask = take_sig_mask;
new.sa_flags = 0;

value.it_interval.tv_sec=0;
value.it_interval.tv_usec=0;
value.it_value.tv_sec=3;
value.it_value.tv_usec=500000;

if(sigaction(SIGVTALRM, &new, &old) == -1){
    perror("can't set signals: ");
    exit(1);
}
setitimer(ITIMER_VIRTUAL, &value, &ovalue);

/** do work to time, take time-out in 3.5 sec */
/** stop clock and get return for measurement */

value.it_value.tv_sec=0;
value.it_value.tv_usec=0;
setitimer(ITIMER_VIRTUAL, &value, &ovalue);
printf("code ran for %d secs and %d micro-secs\n",
    ovalue.it_value.tv_sec, ovalue.it_value.tv_usec);
```


Signal Management (Cont'd)

A non-local go to with signal state

SYNOPSIS

```
#include <setjmp.h>

int sigsetjmp (sigjmp_buf env, int savemask);

void siglongjmp (sigjmp_buf env, int val);
```

DESCRIPTION

These functions are useful for dealing with errors and interrupts encountered in a low-level subroutine of a program. They work as do the `setjmp()` and `longjmp()` routines but if the `savemask` argument is non zero, they also save and restore a process's signal mask upon reactivation of a check point. Since the `sigaction()` call normally blocks further presentation of a signal when in a handler, longjumps out of a handler leave the signal taken blocked until explicitly unblocked by the program. To avoid this, these functions allow the signal mask established during a check point to be restored when that point is reactivated, thereby automatically adjusting the mask for the just taken signal. They are POSIX library routines (3C) from `libc`, NOT system calls.

Interprocess Communication

Since the traditional UNIX process model has been represented with a single thread of execution, and since the system call interface is basically synchronous, many applications are built from more than one process. When multiple process applications are constructed, it is generally necessary to provide a mechanism for the application processes to exchange information. UNIX provides the following IPC support:

- Half-duplex local-host stream communication is supported by **pipes**
 - Un-named pipes have supported IPC in UNIX from the beginning, but they are based on inheritance, and only work well between parent and child or siblings.
 - Named pipes use the name space of the file system to allow any collection of process to connect, but they still are limited by a half-duplex paradigm
- The BSD UNIX Domain Sockets (UDS) provide a full-duplex local-host stream communication facility with either an un-named or named capability.
- The System V IPCs provide named, local-host communication with:
 - A full-duplex datagram message queue facility
 - Shared memory segment support
 - Semaphore synchronization primitives
- The BSD Internet Socket facility provides a full-duplex, named, interhost stream and datagram communication facility

Pipe Based Interprocess Communication

Pipes are fundamentally one-way communication paths which allow one process to write into a pipe and another process to read from a pipe. They support a logical stream of data (generally implemented in the kernel using a message passing STREAMS pseudo driver), with no physical boundaries. A reader must understand how the data in a pipe are organized in order to extract information correctly, since no boundaries are embedded between logical messages.

- Un-named pipes are created using the **pipe()** system call. Typically they are created by a process prior to forking a child, and the resulting allocated pipe file descriptors (two channels, one open for RDONLY and the other for WRONLY) are then inherited by the created child. The child will use either the *write* side channel and leave the *read* side for the parent, or visa-versa.
- Named pipes are created using the **mknod()** system call. Such a pipe is represented by the allocation of a type **p** inode in the file system, which can be opened anytime after creation by any process wishing to communicate using the pipe, and having appropriate credentials to open the object. The **open()** system call used must indicate one of O_RDONLY, O_WRONLY or O_RDWR, and will block the caller until another process calls open() on the named pipe with a compatible mode.

Pipe Based IPC (Cont'd)

IPC with un-named pipe

SYNOPSIS

```
int  pipe (fildes)
int  fildes[2];
```

where:

fildes Address of an array of two
file descriptors; fildes[0]
will hold the smaller channel
number which will be available
for reading only, while the
fildes[1] element will hold the
larger channel which will be
open for writing only.

returns: 0 on success, or -1

Pipe Based IPC (Cont'd)

pipe() cont'd

EXAMPLE:

```
/** A process which talks to itself */
```

```
main()
```

```
{
```

```
    int    parray[2], nread;
```

```
    char   buf[100];
```

```
    if(pipe(parray) == -1){
```

```
        perror("pipe");
```

```
        exit(1);
```

```
    }
```

```
    if(write(parray[1], "hello",6) == -1){
```

```
        perror("write");
```

```
        exit(1);
```

```
    }
```

```
    switch(nread = read(parray[0], buf, sizeof(buf))) {
```

```
        case -1:  perror("read");
```

```
                  exit(1);
```

```
        case 0:  printf("EOF encountered \n");
```

```
                  exit(1);
```

```
        default: printf("read %d bytes: %s \n",
                        nread, buf);
```

```
    }
```

```
}
```

The output:

```
read 6 bytes: hello
```

Using The Standard Tools With Pipes

A one-way communication example:

- The following code assumes a parent process which will carry out some processing on each named object in the current working directory
- Since there is a *naive* UNIX command known as the **ls** command, which will generate the names of all named objects in the current working directory, the application will create a child to run the **ls** command and pipe the results back to the parent application

Using The Standard Tools With Pipes (Cont'd)

EXAMPLE:

```
main()
{
    int pchan[2], pid, done,i;
    char buf[100];

    if(pipe(pchan) == -1){
        perror("pipe");
        exit(1);
    }
    switch( pid = fork() ){
        case -1:  perror("fork");
                 exit(2);
        case  0:  close(1);
                 if( dup(pchan[1] != 1 ){
                     perror("dup");
                     exit(3);
                 }
                 close(pchan[0]);
                 close(pchan[1]);
                 execl( "/bin/ls", "ls", NULL );
                 perror("execl");
                 exit(4);
    }
}
```

Using The Standard Tools With Pipes (Cont'd)

```
/** THE NEXT STATEMENT IS CRITICAL !!! */
close(pchan[1]);
done = 0;
while (1){
    for( i=0; i<100; i++ ){
        if(read(pchan[0], &buf[i], 1) != 1){
            close(pchan[0]);
            done = 1;
            break;                /* BREAK FOR */
        }
        if (buf[i] == '\n'){
            buf[i] = '\0';
            break;                /* BREAK FOR */
        }
    }
    if(done) break;              /* BREAK WHILE */

    /* FILE OBJECT NAME NOW A STRING IN BUF */
    /* ANY PROCESSING WITH NAME DONE HERE */

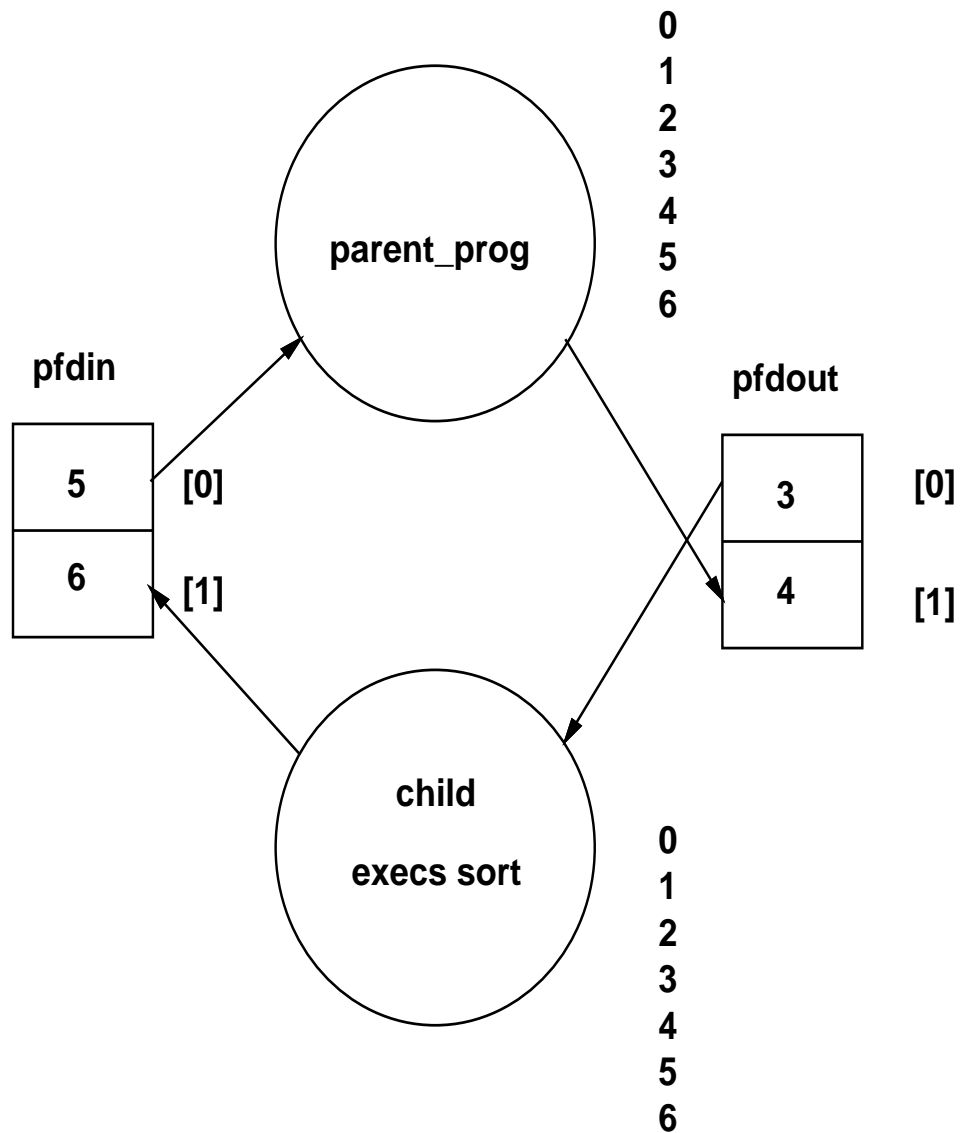
    printf("file object name is %s \n", buf);
}
}
```


Using The Standard Tools With Pipes

A two-way communication example:

- the following example provides a function called `fsort` which provides a process the ability to create a child, load the child with the UNIX **sort** tool, pipe process data to the sort tool for sorting and read back the sorted data
- while the shell does not provide a two way piping syntax, it is quite easy to do from an application. The system programmer must be careful however, since deadlock can be a problem in this situation

A Two-Way Communication Example



Using The Standard Tools With Pipes (Cont'd)

EXAMPLE:

```
#include <stdio.h>
#include <fcntl.h>

main()
{

    int      pfdout[2], pfdin[2], fd, nread;
    char     buf[512];

    if(pipe(pfdout) == -1 || pipe(pfdin) == -1)
    {
        perror("pipe");
        exit(1);
    }

    switch(fork())
    {
        case -1:      perror("fork");
                     exit(2);
```

Using The Standard Tools With Pipes (Cont'd)

```
case 0:    if(close(0) == -1)
           {
               perror("pipe");
               exit(1);
           }
           if(dup(pfdout[0]) != 0)
           {
               perror("dup");
               exit(1);
           }
           if(close(1) == -1)
           {
               perror("pipe");
               exit(1);
           }
           if(dup(pfdin[1]) != 1)
           {
               perror("dup");
               exit(1);
           }
           if(close(pfdout[0]) == -1 ||
              close(pfdout[1]) == -1 ||
              close(pfdin[0]) == -1 ||
              close(pfdin[1]) == -1 )
           {
               perror("close");
               exit(1);
           }
           execlp("grep", "grep", "123", NULL);
           perror("execlp");
           exit(1);
}
```

Using The Standard Tools With Pipes (Cont'd)

```
if(close(pfdout[0]) == -1 || close(pfdin[1]) == -1)
{
    perror("close");
    exit(1);
}
if((fd = open("data", O_RDONLY, 0)) == -1)
{
    perror("open");
    exit(1);
}
while((nread = read(fd, buf, sizeof(buf))) != 0)
{
    if(nread == -1)
    {
        perror("read");
        exit(1);
    }
    if(write(pfdout[1], buf, nread) == -1)
    {
        perror("write");
        exit(1);
    }
}
if(close(fd) == -1 || close(pfdout[1]) == -1)
{
    perror("close");
    exit(1);
}
```

Using The Standard Tools With Pipes (Cont'd)

```
while((nread = read(pfdin[0], buf, sizeof(buf))) != 0)
{
    if(nread == -1)
    {
        perror("read");
        exit(1);
    }
    if(write(1, buf, nread) == -1)
    {
        perror("write");
        exit(1);
    }
}
close(pfdin[0]);
}
```

Pipe Based IPC (Cont'd)

Duplicating file descriptors (channel numbers)

SYNOPSIS

```
int  dup (fildes)
int  fildes;
```

where:

fildes A valid, active file descriptor

SYNOPSIS

```
int  dup2 (old_fildes, new_fildes)
int  old_fildes;
int  new_fildes;
```

where:

old_fildes A valid, active file descriptor

new_fildes Another file descriptor

DESCRIPTION

Dup return the smallest unopen file descriptor and duplicates that descriptor to point to fildes. Dup2 combines the functionality of the dup and close operations. If old_fildes is an active, valid descriptor and new_fildes is a valid descriptor (active or not), new_fildes is made a duplicate of old_fildes. If old_fildes and new_fildes already refer to the same object pointer, no changes occur. In all other situations in which new_fildes is active, it is closed before being made a duplicate of old_fildes. The close-on-exec flag is set so the descriptor remains open across exec(2) operations.

Pipe Based IPC (Cont'd)

Named pipes can be created with the **mknod()** call as previously discussed. There is also a *libc* library convenience routine called *mkfifo()* which call the **mknod()** system call for you:

Make a named pipe

SYNOPSIS

```
#include <sys/types.h>
#include <sys/stat.h>

int mkfifo (const char *path, mode_t mode);
```

DESCRIPTION

The *mkfifo* routine creates a new FIFO special file named by the pathname pointed to by *path*. The mode of the new FIFO is initialized from *mode*. The file permission bits of the mode argument are modified by the process's file creation mask [see *umask(2)*]. The FIFO's owner id is set to the process's effective user id. The FIFO's group id is set to the process's effective group id, or if the *S_ISGID* bit is set in the parent directory then the group id of the FIFO is inherited from the parent. *mkfifo()* calls the system call *mknod()* to make the file.